Supplementary materials: Full code base for simulations and figures

Melissa Chapman[a], Serge Wiltshire[a], Timothy Bowles[a], Liz Carlisle[c], Fredrico Castillo[a], David Gonthier[e], Alastair Iles[a], Daniel Karp[d], Claire Kremen[a,1], Jeff Liebert[f], Elissa Olimpi[d], Matthew Ryan[f], Amber Sciligo[g], Jennifer Thompson[a], Hannah Waterhouse[a], Kenzo Eszquivel[a], Sasha Gennet[h], Carl Boettiger[a,*]

[a]*Dept. of Environmental Science, Policy, and Management, University of California Berkeley, Berkeley, CA, USA*
[b]*Institute of Resources, Environment and Sustainability, University of British Columbia, Vancouver, BC, Canada*
[c]*Environmental Studies, University of California Santa Barbara, Santa Barbara, CA, USA*
[d]*Wildlife, Fish, and Conservation Biology, University of California Davis, Davis, CA, USA*
[e]*Dept. of Entomology, University of Kentucky, Lexington, KY, USA*
[f]*School of Integrated Plant Science, Cornell University, Cornell, NY, USA*
[g]*The Organic Center, Washington, DC, USA*
[h]*The Nature Conservancy, Arlington, VA, USA*

Abstract

This document gives the full codebase used to define and parameterize the DFS-MDP model, run all simulations, and produce the final plots that appear in the accompanying manuscript.

Initializations

Load libraries

```r
rm(list = ls(all.names = TRUE)) # clean environment
# data / MDP libraries
library(tidyverse)
library(Hmisc)
library(dplyr)
library(quantmod) # for findPeaks fxn
library(MDPtoolbox)
# plotting libraries
library(corrplot)
library(RColorBrewer)
library(ggthemes)
library(ggpubr)
library(ggplot2)
library(hrbrthemes)
library(Cairo)
library(extrafont)
extrafont::loadfonts()
ggplot2::theme_set(hrbrthemes::theme_ipsum_rc())
library(patchwork) # devtools::install_github("thomasp85/patchwork")
library(ggtext) # install cli, gh, devtools::install_github("clauswilke/ggtext")
```

*Corresponding Author

Global settings

```
pal <- solarized_pal()(7) # global color pallette
set.seed(4) # random seed
```

Set MDP parameters

```
# define action and state space
actions <- seq(0, 1, length = 100) # vector of possible action values
states <- seq(0, 1.5, length = 100) # vector of possible state values
# set global mdp parameters
params <- list(
  benefit = 1.57, # benefit param for MDP
  cost = 1, # cost param for MDP
  sigma = 0.1, # ecological change stochasticity
  r = 0.1, # ecological change rate
  discount = 0.97, # myopic discount factor
  Tmax = 20 # max simulation time
  )
```

Define state transition probability and utility functions

```
#' Calculate transition probability based on current state and action.
#'
#' @param s The current state.
#' @param a The current action.
#' @param params Must be a list with element r.
#' @return The state transition probability.
transition_fn <- function(s, a, params) s + params$r * (a - s)
#' Calculate utility based on current state and action.
#'
#' @param s The current state.
#' @param a The current action.
#' @param params Must be a list with elements [benefit, cost].
#' @return The state utility.
utility_fn <- function(s, a, params) params$benefit * s - params$cost * a
```

Define MDP model function to calculate P and U matrices

```
#' Calculate transition probability matrix and utility matrix based on MDP parameters.
#'
#' @param states A vector representing the state space.
#' @param actions A vector representing the action space.
#' @param params Must be a list with elements [benefit, cost, sigma].
#' @param transition_fn A function(state, action, params) giving the transition
#' probability to the next state.
#' @param utility_fn A function(state, action, params) giving the value of the next state.
#' @return A list containing the transition probability matrix \code{P} and utility
#' matrix \code{U}.
continuous_model <- function(states, actions, params, transition_fn, utility_fn){
```

```r
    transition_matrix <- function(states, actions, f, params){
      n_s <- length(states)
      n_a <- length(actions)
      transition <- array(0, dim = c(n_s, n_s, n_a))
      for(i in 1:n_a){
        for (k in 1:n_s) {
          nextpop <- transition_fn(states[k], actions[i], params)
          if (nextpop <= 0) {
            x  <- c(1, rep(0, n_s - 1))
          } else {
            x <- truncnorm::dtruncnorm(
              states, 0, max(states), nextpop, params$sigma) # assumes truncated normal error
            if(sum(x) <= 0){
              x  <- c(1, rep(0, n_s - 1))
            } else {
              x <- x / sum(x)
            }
          }
          transition[k, , i] <- x
        }
      }
      if(any(is.na(transition))) stop("error creating transition matrix")
      transition
    }
    utility_matrix <- function(states, actions, utility_fn, params){
      utility <- array(dim=c(length(states), length(actions)))
      for(i in 1:length(states)){
        for(j in 1:length(actions)){
          utility[i,j] <- utility_fn(states[i], actions[j], params)
        }
      }
      utility
    }
    list(P = transition_matrix(states, actions, f, params),
         U = utility_matrix(states, actions, utility_fn, params))
}
```

Define main simulation function

```r
#' Run DFS-MDP simulation for a single agent based on MDP simulation parameters.
#'
#' @param P Transition probability matrix, output from continuous_model().
#' @param U Utility matrix, output from continuous_model().
#' @param policy Optimal decision policy.
#' @param discount Myopic discount factor.
#' @param x0 A vector of initial states.
#' @param Tmax Max time for the simulation.
#' @return A dataframe with time, state, action, and value.
sim_mdp <- function(P, U, policy, discount, x0, Tmax){

  n_states <- dim(P)[1]
```

```
    state <- action <- value <- numeric(Tmax+1)
    state[1] <- x0
    tsteps <- 1:(Tmax+1)
    for(t in tsteps){
      # select action, determine value, transition to next state
      action[t] <- policy[state[t]]
      value[t] <- U[state[t], action[t]] * discount^(t-1)
      state[t+1] <- sample(1:n_states, 1, prob = P[state[t], , action[t]])
    }
    data.frame(time = 0:Tmax, state = state[tsteps],
               action = action[tsteps], value = value[tsteps])
}
```

Set global simulation parameters

```
reps <- 500 # number of replicate simulations to run
init <- truncnorm::rtruncnorm(reps, 0, 1, 0.5, 0.2) %>% # sample init s from norm dist
  map_int(function(x) which.min(abs(x - states))) # map to values in states vector
```

"Base-case" simulation

Define and solve MDP for base-case

```
# compute P and U matrices for the MDP
model <- continuous_model(states, actions, params, transition_fn, utility_fn)
# solve the MDP to find the optimal decision policy
base_soln <- MDPtoolbox::mdp_value_iteration(model$P, model$U, params$discount)
```

```
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
```

```
base_soln_df <- tibble(state = states, action = actions[base_soln$policy])
```

Run base-case simulation

Define timeseries and density plotting functions

```
# timeseries plot of multiple repetitions
sim_plot_ts <- function(sims, title = ggtitle(NULL), tpos = "plot", ytxtoff = FALSE,
                        endcol = pal[1], annotate = FALSE, an_x = NULL, an_lab = NULL,
                        dnmarmod = 0, upmarmod = 0, lmarmod = 0, rmarmod = 0){
  df <- sims %>%
    select(-value) %>% # tidy
    mutate(state = states[state], action = actions[action]) # rescale
  Tmax <- max(sims$time)
  stcol <- col2rgb(endcol)
  stcol <- stcol/5
  stcol <- rgb(t(stcol), maxColorValue=255)
  ytitc <- "black"
  ytxtc = "gray30"
  if (ytxtoff) {
```

```r
    ytitc <- NA
    ytxtc <- NA
  }
  p <- df %>%
    ggplot(aes(time, state, group = reps, col = time)) +
    geom_path(alpha = 0.1, show.legend = FALSE) +
    title +
    labs(x="Decision cycle", y="Ecosystem service state") +
    scale_x_continuous(limits = c(0,Tmax), breaks = seq(0, Tmax, by=5),
                       expand = c(.01,.01)) +
    scale_y_continuous(limits = c(0,1.5), expand = c(.02,.02)) +
    scale_color_gradient(low=stcol, high=endcol) +
    theme(axis.text.x=element_text(size=10),
          axis.text.y=element_text(size=10, color = ytxtc),
          axis.title.x=element_text(size=10),
          axis.title.y=element_text(size=10, color = ytitc),
          plot.title = element_text(size = 10, face = "bold"),
          plot.title.position = tpos,
          panel.grid.minor = element_blank(),
          plot.margin=grid::unit(c(5+upmarmod,5+rmarmod,5+dnmarmod,5+lmarmod), "mm"))
  if(annotate) {
    p <- p +
      geom_vline(xintercept = an_x, linetype="dashed", color = "red3", size=.5) +
      annotate('label', x = an_x + 2, y = 1.3, label = an_lab, hjust = 0, vjust = .5,
               family = "Roboto", size = 3.25, label.padding = unit(.15, "lines"),
               label.size = 0, alpha = .6) +
      annotate("segment", x = an_x + 1.75, xend = an_x + .5, y = 1.3, yend = 1.3,
               size=.5, arrow=arrow(length = unit(0.22, "cm")))
  }
  p
}
# density plot showing initial ES distribution and final distribution
sim_plot_dens <- function(sims, title = ggtitle(NULL), tpos = "plot", endcol = pal[1],
                          lab_lo_peak = FALSE, lab_hi_peak = FALSE,
                          dnmarmod = 0, upmarmod = 0, lmarmod = 0, rmarmod = 0){
  df <- sims %>%
    mutate(state = states[state]) %>% # rescale
    select(state, time)
  Tmax <- max(sims$time)
  stcol <- col2rgb(endcol)
  stcol <- stcol/5
  stcol <- rgb(t(stcol), maxColorValue=255)
  p <- df %>% filter(time %in% c(0, Tmax))  %>%
    ggplot() + geom_density(aes(state, group = time, fill = time, color = time),
                            alpha=0.8) +
    coord_flip() +
    title +
    labs(x="", y="Density", fill="Time") +
    scale_x_continuous(limits = c(0,1.5), expand = c(.02,.02)) +
    scale_y_continuous(breaks = scales::pretty_breaks(n = 2)) +
    scale_fill_gradient(low=stcol, high=endcol, guide = guide_colorbar(barwidth = .5),
```

```
                              breaks=c(0, Tmax)) +
     scale_color_gradient(low=stcol, high=endcol, guide = NULL) +
     theme(axis.text.x=element_text(size=10),
           axis.text.y=element_blank(),
           axis.title.x=element_text(size=10),
           axis.title.y=element_text(size=10),
           legend.text = element_text(size=10),
           legend.title = element_text(size=10),
           legend.box.margin=margin(0,0,0,-5),
           plot.title = element_text(size = 10, face = "bold"),
           plot.title.position = tpos,
           panel.grid.minor = element_blank(),
           plot.margin=grid::unit(c(5+upmarmod,5+rmarmod,5+dnmarmod,5+lmarmod), "mm"))
 if(lab_lo_peak) {
   ymax <- ggplot_build(p)$layout$panel_scales_y[[1]]$range$range[2]
   peak_lo <- ggplot_build(p)$data[[1]] %>%
     filter(group == 2, x < .4) %>%
     arrange(desc(y)) %>%
     select(x,y) %>%
     filter(row_number()==1)
   # dotted line
   p <- p +
     annotate('segment', x = peak_lo$x, xend = peak_lo$x, y = peak_lo$y - .1 * ymax,
              yend = peak_lo$y + .1 * ymax, linetype="dotted")
   # text annotation (decide whether to place left or right of line)
   if(peak_lo$y > .7 * ymax) {
     p <- p +
       annotate('label', x = peak_lo$x, y = peak_lo$y - .32 * ymax,
                label = sprintf('%.2f', peak_lo$x), hjust = .5, vjust = .5,
                family = "Roboto Condensed", size = 3,
                label.padding = unit(.15, "lines"), label.size = 0, alpha = .8)
   } else {
     p <- p +
       annotate('label', x = peak_lo$x, y = peak_lo$y + .26 * ymax,
                label = sprintf('%.2f', peak_lo$x), hjust = .5, vjust = .5,
                family = "Roboto Condensed", size = 3,
                label.padding = unit(.15, "lines"), label.size = 0, alpha = .8)
   }
 }
 if(lab_hi_peak){
   ymax <- ggplot_build(p)$layout$panel_scales_y[[1]]$range$range[2]
   peak_hi <- ggplot_build(p)$data[[1]] %>%
     filter(group == 2, x > .6) %>%
     arrange(desc(y)) %>%
     select(x,y) %>%
     filter(row_number()==1)
   # dotted line
   p <- p +
     annotate('segment', x = peak_hi$x, xend = peak_hi$x, y = peak_hi$y - .1 * ymax,
              yend = peak_hi$y + .1 * ymax, linetype="dotted")
   # text annotation (decide whether to place left or right of line)
```

```
    if(peak_hi$y > .7 * ymax) {
      p <- p +
        annotate('label', x = peak_hi$x, y = peak_hi$y - .32 * ymax,
                 label = sprintf('%.2f', peak_hi$x), hjust = .5, vjust = .5,
                 family = "Roboto Condensed", size = 3,
                 label.padding = unit(.15, "lines"), label.size = 0, alpha = .8)
    } else {
      p <- p +
        annotate('label', x = peak_hi$x, y = peak_hi$y + .26 * ymax,
                 label = sprintf('%.2f', peak_hi$x), hjust = .5, vjust = .5,
                 family = "Roboto Condensed", size = 3,
                 label.padding = unit(.15, "lines"), label.size = 0, alpha = .8)
    }
  }
  p
}
```

Generate base-case simulation figure

```
ggarrange(
  sim_plot_ts(sims_baseline, title = ggtitle("A. State timeseries"), rmarmod = -5),
  sim_plot_dens(sims_baseline, title = ggtitle("B. Initial and final state"),
                tpos = "panel", lab_lo_peak = TRUE, lab_hi_peak = TRUE, lmarmod = -5),
  widths = c(2.75, 2), nrow = 1, ncol = 2
)
```
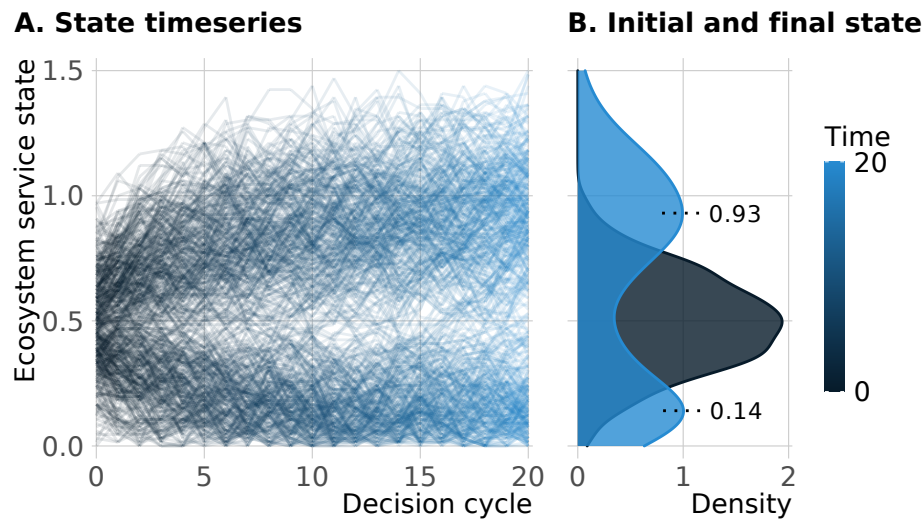

Figure 1: Initial ecosystem states are distributed normally (mean = 0.5; S.D. = 0.2; truncated at [0,1]). Agents follow decision strategy π as shown in Fig 2 until $t = 20$. (A) Ecosystem state of each agent over time (500 simulations). (B) Initial ES distribution (dark blue) and final bimodal distribution at $t = 20$ (light blue).

Define base-base optimal decision strategy plotting function

7

```
# solution point plot with threshold annotated
soln_plot <- function(soln_df, tpt) {
  ggplot(soln_df, aes(state,action)) +
    geom_point(size = 1) +
    geom_vline(xintercept = tpt, linetype="dashed", color = "red3", size=.7) +
    annotate('label', x = tpt + .15, y = .375, label = "Tipping point",
             hjust = 0, vjust = .5, family = "Roboto", size = 3.25,
             label.padding = unit(.15, "lines"), label.size = 0, alpha = .65) +
    annotate("segment", x = tpt + .15, xend = tpt + .025, y = .375, yend = .375,
             size=.5, arrow=arrow(length = unit(0.22, "cm"))) +
    labs(x="Ecosystem service state", y="Optimal adoption") +
    scale_x_continuous(limits = c(0,NA), expand = c(.01,.01)) +
    scale_y_continuous(limits = c(0,1), expand = c(.01,.01),
                       labels = scales::percent_format(accuracy = 1)) +
    theme(axis.text.x=element_text(size=10),
          axis.text.y=element_text(size=10),
          axis.title.x=element_text(size=10),
          axis.title.y=element_text(size=10),
          panel.grid.minor = element_blank(),
          plot.margin=grid::unit(c(5,5,5,5), "mm"))
}
```

Generate optimal policy plot

```
state99th <- sims_baseline %>%
  mutate(state = states[state]) %>%
  pull(state) %>%
  quantile(.99) %>%
  unname()
tpt <- 0
for(si in 1:length(states)-1) {
  if(base_soln_df$action[si] < .1 && base_soln_df$action[si+1] > .9) {
    tpt <- base_soln_df$state[si] + ((base_soln_df$state[2] - base_soln_df$state[1]) / 2)
  }
}
soln_plot(base_soln_df %>% filter(state <= state99th), tpt = tpt)
```

Figure 2: Optimal decision strategy π as a function of ES state, showing a tipping point at $s \approx 0.48$. The upper x axis limit is the 99th percentile of observed states in our simulation results ($s \approx 1.3$).

Land tenure experiment

Calculate short-tenure MDP solution

Define density curve comparison plotting function

```r
# plot formatting helper functions
# increase vertical spacing between legend keys
# written by @clauswilke
draw_key_polygon3 <- function(data, params, size) {
  lwd <- min(data$size, min(size) / 4)
  grid::rectGrob(
    width = grid::unit(0.3, "npc"),
    height = grid::unit(0.8, "npc"),
    gp = grid::gpar(
      col = data$colour,
      fill = alpha(data$fill, data$alpha),
      lty = data$linetype,
      lwd = lwd * .pt,
      linejoin = "mitre"
    ))
}
# register new key drawing function,
# the effect is global & persistent throughout the R session
GeomDensity$draw_key = draw_key_polygon3
# plot comparing final and initial density curves of two simulations
sim_plot_dens_comp <- function(sims1, sims2, sims_base = "init",
                               label1 = "Gp. A Final", label2 = "Gp. B Final",
                               label_base = "Initial", title = ggtitle(NULL),
                               tpos = "plot", cvec = c(pal[2], pal[1], pal[3]),
                               dnmarmod = 0, upmarmod = 0, lmarmod = 0, rmarmod = 0){
  if (sims_base == "init") {
```

9

```r
    df_base <- sims1 %>%
      mutate(state = states[state]) %>% # rescale
      filter(time %in% 0) %>%
      select(state) %>%
      add_column(id = label_base)
  } else {
    Tmax_base <- max(sims_base$time)
    df_base <- sims_base %>%
      mutate(state = states[state]) %>% # rescale
      filter(time %in% Tmax_base) %>%
      select(state) %>%
      add_column(id = label_base)
  }
  Tmax1 <- max(sims1$time)
  df1 <- sims1 %>%
    mutate(state = states[state]) %>% # rescale
    filter(time %in% Tmax1) %>%
    select(state) %>%
    add_column(id = label1)
  Tmax2 <- max(sims2$time)
  df2 <- sims2 %>%
    mutate(state = states[state]) %>% # rescale
    filter(time %in% Tmax1) %>%
    select(state) %>%
    add_column(id = label2)
  rbind(df_base, df1, df2) %>% mutate(id = factor(id, unique(id))) %>%
    ggplot() + geom_density(aes(state, group = id, fill = id, color = id), alpha=0.4) +
    coord_flip() +
    title +
    labs(x="", y="Density") +
    scale_x_continuous(limits = c(0,1.5), expand = c(.02,.02)) +
    scale_y_continuous(breaks = scales::pretty_breaks(n = 2)) +
    scale_fill_manual(values = cvec) +
    scale_color_manual(values = cvec) +
    theme(axis.text.x=element_text(size=10),
          axis.text.y=element_blank(),
          axis.title.x=element_text(size=10),
          axis.title.y=element_text(size=10),
          legend.text=element_text(size=9),
          legend.title = element_blank(),
          legend.key.size = unit(12, "mm"),
          legend.spacing.x = unit(-3, 'mm'),
          legend.box.margin=margin(0,0,0,-18),
          plot.title = element_text(size = 10, face = "bold"),
          plot.title.position = tpos,
          panel.grid.minor = element_blank(),
          plot.margin=grid::unit(c(5+upmarmod,5+rmarmod,5+dnmarmod,5+lmarmod), "mm"))
}
```

Generate land tenure experiment plot

```r
ggarrange(
  sim_plot_ts(sims_short_tenure,
              title = ggtitle("A. Short-tenure state timeseries"),
              endcol = pal[3], rmarmod = -5),
  sim_plot_dens_comp(sims_baseline, sims_short_tenure,
                     label1 = "Long\ntenure\nfinal",
                     label2 = "Short\ntenure\nfinal",
                     title = ggtitle("B. State distribution"),
                     tpos = "panel", lmarmod = -5),
  widths = c(2.75, 2), nrow = 1, ncol = 2
)
```
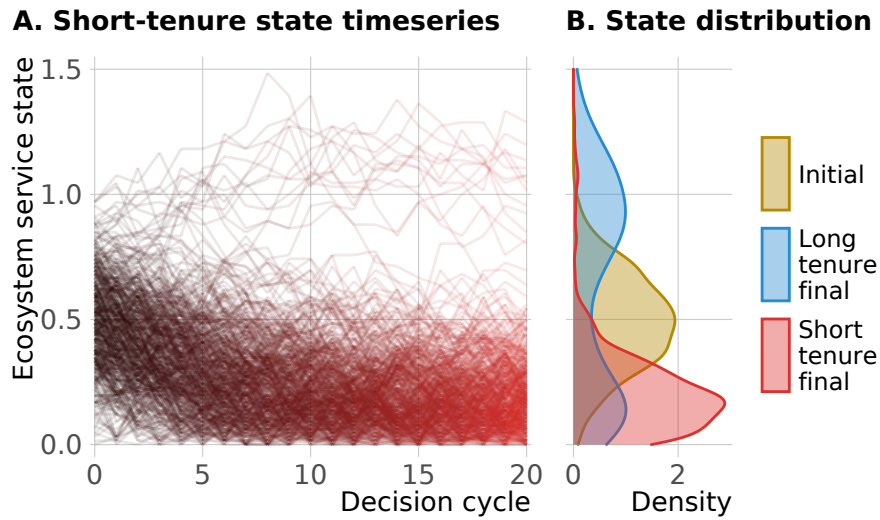

Figure 3: The simulation is identical to that in Fig 1, but the MDP is solved under a finite, 20-year time horizon. (A) Result of short land tenure on ES state over time. (B) Comparison between final state distribution of short- vs. long-tenure model runs.

Incentive duration experiment

Incentive simulation parameters and setup

```r
# set incentive simulation parameters
T_incent_1 <- 2
T_incent_2 <- 10
# calc modified cost during incentive = 1 - 2 / T_incent
C_incent_1 <- 1 - 2 / T_incent_1
C_incent_2 <- 1 - 2 / T_incent_2
# incentive simulation fxn
incent_sim <- function(params) {
  incent_model <- continuous_model(states, actions, params,
                                   transition_fn, utility_fn)
  incent_soln <- MDPtoolbox::mdp_value_iteration(incent_model$P,
                                                 incent_model$U, params$discount)
```

```
  # simulate incentive period
  start <- sim(incent_soln, Tmax = params$T_incent, x0 = init)
  xi <- start %>% filter(time == params$T_incent) %>% pull(state)
  # pad to Tmax with baseline decision rule
  rest <- sim(base_soln, Tmax = params$Tmax - params$T_incent, x0 = xi) %>%
    filter(time!=0) %>% mutate(time = time + params$T_incent)
  bind_rows(start, rest)
}
```

Run incentives simulation

```
# run incentive simulation
sims_incent_abrupt <- incent_sim(list(benefit = params$benefit,
                                  cost = C_incent_1,
                                  sigma = params$sigma,
                                  r = params$r,
                                  discount = params$discount,
                                  T_incent = T_incent_1, Tmax = params$Tmax))
```

```
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
```

```
sims_incent_sust <- incent_sim(list(benefit = params$benefit,
                                cost = C_incent_2,
                                sigma = params$sigma,
                                r = params$r,
                                discount = params$discount,
                                T_incent = T_incent_2, Tmax = params$Tmax))
```

```
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
```

Plot incentive simulation

```
ggarrange(
  sim_plot_ts(sims_incent_abrupt, title = ggtitle("A. Concentrated incentive"),
            annotate = TRUE, an_x = T_incent_1,
            an_lab = sprintf("%i cycles\n%i%% off", T_incent_1,
                            round((1-C_incent_1)*100)),
            rmarmod = -10, endcol = pal[5]),
  sim_plot_ts(sims_incent_sust, title = ggtitle("B. Sustained incentive"),
            annotate = TRUE, an_x = T_incent_2,
            an_lab = sprintf("%i cycles\n%i%% off", T_incent_2,
                            round((1-C_incent_2)*100)),
            tpos = "panel", ytxtoff = T, endcol = pal[6], rmarmod = -5, lmarmod = -5),
  sim_plot_dens_comp(sims_incent_abrupt, sims_incent_sust, sims_base = sims_baseline,
                  label1 = "Concentrated \nincentive", label2 = "Sustained\nincentive",
                  label_base = "No\nincentive",
                  title = ggtitle("C. Final state distribution"), tpos = "panel",
                  cvec = c(pal[2], pal[5], pal[6]), lmarmod = -5, rmarmod = -5),
  nrow = 1, ncol = 3, widths = c(1,1,1)
)
```
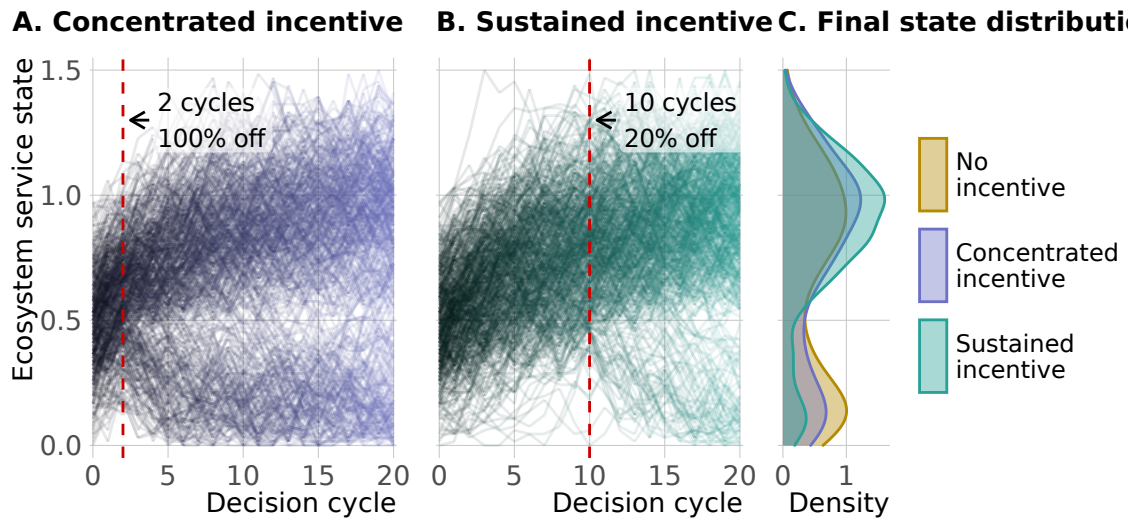
Figure 4: Starting from the same initial states as Fig 1, ES state timeseries are shown for (A) a large, abrupt incentive (100% of adoption expenses are covered for two years) vs. (B) a smaller, more sustained incentive (adoption cost is 80% of baseline for 10 years). Ignoring discounting, both packages have the same total cost to the funder (the equivalent of 2 years' worth of full adoption cost offsets). After the incentive period, agents adjust their decision rules to that of the base case (i.e. no incentive) until $t = 20$. (C) Shows that the sustained incentive ultimately drove more DP adoption.

Temporal dynamics experiment

Define "benefit sweep" simulation function

```
sim_b_sweep <- function(scenarios, inf_fin, t_horiz = 2) {

  sim_b_sweep_scenario <- function(params) {
    model <- continuous_model(states, actions, params, transition_fn, utility_fn)
    if(inf_fin == "infinite") {
      soln <- MDPtoolbox::mdp_value_iteration(model$P, model$U, params$discount)
    } else if(inf_fin == "finite") {
      soln <- MDPtoolbox::mdp_finite_horizon(model$P, model$U, params$discount,
                                              N = t_horiz)
    } else stop('inf_fin must be "infinite" or "finite"')
    sim(soln, Tmax = params$Tmax) %>%
      filter(time == params$Tmax)
  }

  scenarios %>%
    purrr::transpose() %>%
    furrr::future_map_dfr(sim_b_sweep_scenario, .id = "id") %>%
    left_join(scenarios, by = "id") %>%
    mutate(state = states[state])
}
```

Parameterize and run "benefit-sweep" simulation scenarios

13

```r
chans_scenarios <-
  tidyr::crossing(cbr = seq(.55, .7, length.out = 40),
                  cost = params$cost,
                  sigma = params$sigma,
                  r = params$r,
                  discount = params$discount,
                  Tmax = params$Tmax) %>%
  tibble::rownames_to_column("id") %>%
  mutate(benefit = cost / cbr)
chans_sims <- sim_b_sweep(chans_scenarios, "infinite")
```

```
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
```

```
fast_r_scenarios <-
  tidyr::crossing(cbr = seq(.825, .975, length.out = 40),
                  cost = params$cost,
                  sigma = params$sigma,
                  r = 0.95,
                  discount = params$discount,
                  Tmax = params$Tmax) %>%
  tibble::rownames_to_column("id") %>%
  mutate(benefit = cost / cbr)
fast_r_sims <- sim_b_sweep(fast_r_scenarios, "infinite")
```

```
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
## [1] "MDP Toolbox: iterations stopped, epsilon-optimal policy found"
```

```r
myopic_d_scenarios <-
  tidyr::crossing(cbr = seq(.025, .175, length.out = 40),
                  cost = params$cost,
                  sigma = params$sigma,
                  r = params$r,
                  discount = params$discount,
                  Tmax = params$Tmax) %>%
  tibble::rownames_to_column("id") %>%
  mutate(benefit = cost / cbr)
myopic_d_sims <- sim_b_sweep(myopic_d_scenarios, "finite", t_horiz = 2)
```

Define "benefit sweep" simulation helper and plotting functions

```r
# find peaks fxn for benefit sweep experiment
b_sweep_peaks <- function(sims, smoothing = 1.5,
                          ythresh = 0.125, output_bizone = FALSE) {

  bimin = 0
  bimax = 0
  peaks <- tibble(cbr = numeric(), peak = numeric())
  for(this_cbr in unique(sims$cbr)) {
    dens <- density((sims %>% filter(cbr == this_cbr))$state, adjust = smoothing)
    dens <- tibble(x = dens$x, y = dens$y) %>% filter(y >= ythresh)
    ps <- dens$x[findPeaks(dens$y)]
    for(p in ps) {
      peaks <- add_row(peaks, cbr = this_cbr, peak = p)
    }
    if(output_bizone && length(ps) > 1) {
      if(bimin == 0) {
        bimin = this_cbr
      }
      bimax = this_cbr
    }
  }
  if(output_bizone) {
    c(bimin, bimax)
  } else {
    peaks
  }
}
# plot benefit sweep experiment results
plt_sim_b_sweep <- function(sims, title = ggtitle(NULL), ylab = "",
                            tpos = "plot", col = pal[1],
                            upmarmod = 0, lmarmod = 0, dnmarmod = 0, rmarmod = 0) {
  dkcol <- col2rgb(col)
  dkcol <- dkcol/5
  dkcol <- rgb(t(dkcol), maxColorValue=255)
  sims %>%
    ggplot(aes(x = state, group = cbr, color = cbr)) +
    geom_line(stat='density', size=.5, alpha=.4) +
    title +
```

```r
    labs(x = "ES state", y = ylab, color = "c:b") +
    scale_x_continuous(limits = c(0, 1.5),
                       breaks = scales::pretty_breaks(n = 3), expand = c(.01,.01)) +
    scale_y_continuous(breaks = scales::pretty_breaks(n = 2), expand = c(.01,.01)) +
    scale_color_gradient(low = dkcol, high = col, guide = guide_colorbar(barwidth = .5),
                         breaks=c(min(sims$cbr), max(sims$cbr))) +
    theme(axis.text.x=element_text(size=10),
          axis.text.y=element_text(size=10),
          axis.title.x=element_text(size=10),
          axis.title.y=element_text(size=10),
          legend.text = element_text(size=10),
          legend.title = element_text(size=10),
          legend.box.margin=margin(0,0,0,-5),
          plot.title = element_text(size = 10, face = "bold"),
          plot.title.position = tpos,
          panel.grid.minor = element_blank(),
          plot.margin=grid::unit(c(5+upmarmod,5+rmarmod,5+dnmarmod,5+lmarmod), "mm"))
}
# plot benefit sweep experiment peaks by cbr
plt_b_sweep_peaks <- function(peaks, bizone = NULL, title = ggtitle(NULL), ylab = "",
                              tpos = "plot", col = pal[1],
                              upmarmod = 0, lmarmod = 0, dnmarmod = 0, rmarmod = 0) {
  dkcol <- col2rgb(col)
  dkcol <- dkcol/5
  dkcol <- rgb(t(dkcol), maxColorValue=255)
  p <- peaks %>%
    ggplot(aes(x = peak, y = cbr, color = cbr)) +
    geom_point(size=1) +
    title +
    labs(x = "state density peak(s)", y = ylab) +
    scale_x_continuous(limits = c(0, 1.5),
                       breaks = scales::pretty_breaks(n = 3), expand = c(.01,.01)) +
    scale_y_continuous(limits = c(min(peaks$cbr), max(peaks$cbr)),
                       breaks = scales::pretty_breaks(n = 3), expand = c(.01,.01)) +
    scale_color_gradient(low = dkcol, high = col, guide = NULL) +
    theme(axis.text.x=element_text(size=10),
          axis.text.y=element_text(size=10),
          axis.title.x=element_text(size=10),
          axis.title.y=element_text(size=10),
          plot.title = element_text(size = 10, face = "bold"),
          plot.title.position = tpos,
          panel.grid.minor = element_blank(),
          plot.margin=grid::unit(c(5+upmarmod,5+rmarmod,5+dnmarmod,5+lmarmod), "mm"))
  if(!is.null(bizone)) {
    p <- p +
      geom_hline(yintercept = bizone[1], linetype="dashed", color = "red3", size=.5) +
      geom_hline(yintercept = bizone[2], linetype="dashed", color = "red3", size=.5) +
      annotate('label', x = .55, y = mean(bizone), label = "Bimodal\nregion",
               hjust = .5, vjust = .5, family = "Roboto", size = 3,
               label.padding = unit(.15, "lines"), label.size = 0, alpha = .65)
  }
```

```
  p
}
```

Generate "benefit sweep" simulation plot

```
ggarrange(
  ggarrange(
    plt_sim_b_sweep(chans_sims,
                    title = ggtitle("A. Forward-looking decision-maker
                                    \n    and adaptive ecological system"),
                    ylab = "Density", col = pal[1], rmarmod = -7, dnmarmod = -2.75) +
      theme(legend.position="none"),
    plt_b_sweep_peaks(b_sweep_peaks(chans_sims),
                      bizone = b_sweep_peaks(chans_sims, output_bizone = TRUE),
                      ylab = "Cost/benefit ratio", col = pal[1],
                      rmarmod = -7, upmarmod = -2.75),
    nrow = 2, ncol = 1, heights = c(1.1,1), align = "v"
  ),
  ggarrange(
    plt_sim_b_sweep(myopic_d_sims,
                    title = ggtitle("B. Short-term
                                    \n    decision strategy"),
                    tpos = "panel", col = pal[4],
                    lmarmod = -3.5, rmarmod = -3.5, dnmarmod = -2.75) +
      theme(legend.position="none"),
    plt_b_sweep_peaks(b_sweep_peaks(myopic_d_sims),
                      col = pal[4], lmarmod = -3.5, rmarmod = -3.5, upmarmod = -2.75),
    nrow = 2, ncol = 1, heights = c(1.1,1), align = "v"
  ),
  ggarrange(
    plt_sim_b_sweep(fast_r_sims,
                    title = ggtitle("C. fast
                                    \n    ecological change rate"),
                    tpos = "panel", col = pal[7], lmarmod = -7, dnmarmod = -2.75) +
      theme(legend.position="none"),
    plt_b_sweep_peaks(b_sweep_peaks(fast_r_sims),
                      col = pal[7], lmarmod = -7, upmarmod = -2.75),
    nrow = 2, ncol = 1, heights = c(1.1,1), align = "v"
  ),
  nrow = 1, ncol = 3, widths = c(1,1,1)
)
```

Figure 5: For three scenarios (coupled human/natural system, overly-myopic decision maker, and overly-fast ecological change), cost/benefit ratio was varied incrementally over 40 values, indicated by color shade, across a $c : b$ range of width 0.15, encompassing the transition between a "never invest" to an "always invest" policy. For each $c : b$, 500 replicate simulations were conducted as in Fig 1. Upper plots show distribution of ES state at $t = 20$ for each $c : b$. Lower plots show density curve peak(s). (A) By coupling a forward-looking decision-maker and a slowly-adapting environment, complex dynamics like alternate stable states can emerge. However, with (B) a short-term decision strategy (solving the MDP over a 2-year time horizon), or (C) a fast ecological change rate ($r = 0.95$), no bimodality is observed.